

IRTM_Project_TudorAndrei_Dumitrascu

January 23, 2022

0.1 IRTM Project 2

Dumitrascu Tudor Andrei

```
[ ]: import spacy
import pandas as pd
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegressionCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from tqdm import tqdm
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.metrics import classification_report, f1_score
from statistics import mean
import numpy as np
```

```
[ ]: # !python -m spacy download en_core_web_sm
```

```
[ ]: nlp = spacy.load("en_core_web_sm", exclude=['ner', 'lemmatizer'])
```

1 Load data

```
[ ]: data = pd.read_csv("Lyrics-Genre-Train.csv").drop(['Song', 'Song year',
↳ 'Artist', 'Track_id'], axis=1)
data_test = pd.read_csv("Lyrics-Genre-Test-GroundTruth.csv").drop(['Song',
↳ 'Song year', 'Artist', 'Track_id'], axis=1)
```

```
[ ]: def generate_pos(x) -> str:
    doc = nlp(x)
    results = []
    for i, token in enumerate(doc):
        results.append(token.pos_)

    return " ".join(results)
```

```
[ ]: data['POS'] = data['Lyrics'].apply(generate_pos)
data_test['POS'] = data_test['Lyrics'].apply(generate_pos)
```

```
[ ]: def train_eval(model, X_train, y_train, X_test, y_test, encoder):
    print(str(model))
    model.fit(X_train, y_train)
    print("Train")
    y_hat = model.predict(X_train)
    print(f1_score(y_train, y_hat, average='weighted'))
    print("Test")
    y_hat = model.predict(X_test)
    print(f1_score(y_test, y_hat, average='weighted'))
```

2 Implementation

2.1 Verse Features

```
[ ]: def pos_count(x, pos):
    return sum([y == pos for y in x.split() ])

# count the part of speech tags in the whole song
```

```
[ ]: poeses = set(data.loc[0, 'POS'].split())
poeses
```

```
[ ]: {'ADJ',
      'ADP',
      'ADV',
      'AUX',
      'CCONJ',
      'DET',
      'NOUN',
      'PRON',
      'PROPN',
      'PUNCT',
      'SPACE',
      'VERB'}
```

```
[ ]: for pos in tqdm(poeses):
    data[ pos + "_count"] = data['POS'].apply(lambda x: pos_count(x, pos))
    data_test[ pos + "_count"] = data_test['POS'].apply(lambda x: pos_count(x,
↪pos))
```

100%| | 12/12 [00:07<00:00, 1.52it/s]

```
[ ]: # Count the number of verses
data['newln_count'] = data['Lyrics'].apply(lambda x: x.count("\n"))
data_test['newln_count'] = data_test['Lyrics'].apply(lambda x: x.count("\n"))
```

```

# Count the avg verse length
data['mean_verse_len'] = data['Lyrics'].apply(lambda x: mean([len(y) for y in x.
    ↪split("\n")]))
data_test['mean_verse_len'] = data_test['Lyrics'].apply(lambda x: mean([len(y)
    ↪for y in x.split("\n")]))
# count the avg no. of words per verse
data['mean_word_count_per_verse'] = data['Lyrics'].apply(lambda x: mean([len(y.
    ↪split()) for y in x.split("\n")]))
data_test['mean_word_count_per_verse'] = data_test['Lyrics'].apply(lambda x:
    ↪mean([len(y.split()) for y in x.split("\n")]))

```

```

[ ]: X_train = data.drop(['Genre', 'Lyrics', 'POS'], axis=1)
X_test = data_test.drop(['Genre', 'Lyrics', 'POS'], axis=1)
enc = LabelEncoder()
y_train = enc.fit_transform(data['Genre'])
y_test = enc.transform(data_test['Genre'])

```

```

[ ]: train_eval(SVC(class_weight='balanced'), X_train, y_train, X_test, y_test, enc)

```

```

SVC(class_weight='balanced')
Train
0.2917416059978742
Test
0.277862611528458

```

```

[ ]: train_eval(DecisionTreeClassifier(class_weight="balanced", criterion='entropy',
    ↪max_depth=10), X_train, y_train, X_test, y_test, enc)

```

```

DecisionTreeClassifier(class_weight='balanced', criterion='entropy',
                        max_depth=10)
Train
0.3493602597226395
Test
0.23033025647020453

```

```

[ ]: train_eval(RandomForestClassifier(n_jobs=-1), X_train, y_train, X_test, y_test,
    ↪enc)

```

```

RandomForestClassifier(n_jobs=-1)
Train
0.9998919615100562
Test
0.29142742830782753

```

2.2 TFIDF

```
[ ]: tfidf = TfidfVectorizer(ngram_range=(1,3), max_features=2000, max_df=0.8,  
    ↪min_df=0.2)  
X_train = tfidf.fit_transform(data['Lyrics'])  
X_test = tfidf.transform(data_test['Lyrics'])  
  
[ ]: enc = LabelEncoder()  
y_train = enc.fit_transform(data['Genre'])  
y_test = enc.transform(data_test['Genre'])  
  
[ ]: train_eval(SVC(class_weight='balanced'), X_train, y_train, X_test, y_test, enc)  
  
SVC(class_weight='balanced')  
Train  
0.5939188926029967  
Test  
0.30550260205341334  
  
[ ]: train_eval(RandomForestClassifier(n_jobs=-1), X_train, y_train, X_test, y_test,  
    ↪enc)  
  
RandomForestClassifier(n_jobs=-1)  
Train  
0.9992441696631217  
Test  
0.295679015623803
```

2.3 Part of speech tag

```
[ ]: tfidf = TfidfVectorizer(ngram_range=(1,3), max_features=300)  
X_train = tfidf.fit_transform(data['POS'])  
X_test = tfidf.transform(data_test['POS'])  
  
[ ]: enc = LabelEncoder()  
y_train = enc.fit_transform(data['Genre'])  
y_test = enc.transform(data_test['Genre'])  
  
[ ]: train_eval(SVC(class_weight='balanced'), X_train, y_train, X_test, y_test, enc)  
  
SVC(class_weight='balanced')  
Train  
0.4188555378773796  
Test  
0.3203567772402991  
  
[ ]: train_eval(RandomForestClassifier(n_jobs=-1), X_train, y_train, X_test, y_test,  
    ↪enc)
```

```

RandomForestClassifier(n_jobs=-1)
Train
0.9998379632304489
Test
0.30130022882875435

```

2.4 Word len

```

[ ]: enc = LabelEncoder()
y_train = enc.fit_transform(data['Genre'])
y_test = enc.transform(data_test['Genre'])

```

```

[ ]: def word_len(string, max_pad=200):
    doc = nlp.tokenizer(string)
    results = []
    for i, token in enumerate(doc):
        results.append(len(token))
    results = np.asarray(results)
    if len(results) < max_pad:
        results = np.pad(results, (0, max_pad - len(results)))
    else:
        results = results[:200]
    return results

```

```

[ ]: X_train = data['Lyrics'].apply(word_len)
X_train = np.stack(X_train.values)

```

```

[ ]: X_test = data_test['Lyrics'].apply(word_len)
X_test = np.stack(X_test.values)

```

```

[ ]: min_max = MinMaxScaler().fit(X_train)
X_train = min_max.transform(X_train)
X_test = min_max.transform(X_test)

```

```

[ ]: train_eval(SVC(class_weight='balanced'), X_train, y_train, X_test, y_test, enc)

```

```

SVC(class_weight='balanced')
Train
0.5391034891259431
Test
0.1806859460442372

```

```

[ ]: train_eval(RandomForestClassifier(n_jobs=-1), X_train, y_train, X_test, y_test,
    ↪ enc)

```

```

RandomForestClassifier(n_jobs=-1)
Train
0.9998919615100562

```

Test

0.16504548553924356

[]: